

Soul Protocol Release Notes

On Memory, Rhythm, and Local Intelligence

Lewis Chen · Dragon Boat Festival 2026

Why I Built This

Three months ago, shortly after joining the lab, I stumbled upon OpenClaw. Its architecture hit me with incredible force—for the first time, within a purely silicon-based computing system, I felt something resembling the texture of a "life cycle."

I shared this feeling with friends and colleagues, but most found it hard to relate. Perhaps to them, the idea seemed too far removed from convention.

Not being understood didn't matter. I decided to verify it myself. Working in my spare time with the tools available today, I built Soul Protocol piece by piece.

The current version's code isn't exactly elegant—there's plenty of redundancy. But it already meets basic practical needs: one-command deployment, ready to use out of the box. At least in terms of architectural interpretability, overall design philosophy, and deployment convenience, it has achieved its milestone goals; what comes next is just continuous refinement and bug fixing through real-world use.

Choosing to open-source it during the Dragon Boat Festival holiday is simply my way of putting these thoughts out there, so more people can experience them firsthand.

The Underlying Logic of the Architecture

If LLMs replicate human thinking patterns on computers, then what this plugin aims to do is replicate human memory mechanisms at the Agent level—and further out, the ultimate direction of this project is to explore the engineering replication of human intelligence.

So the entire architecture starts from reverse-engineering human memory and rhythm:

What would a person do in a corresponding scenario? What are human behavioral patterns? And if we want to functionally replicate such behavior, how would we implement it in code?

Following this line of thinking came the layered memory architecture, which achieves basic context management engineering through layered injection; came the Identity

Layer Constitution, used to anchor the AI's self-cognition; and came the Freudian-like tripartite division of the L3 layer, used to construct the underlying shape of personality and facts.

OpenClaw's dreaming design in April surprised me greatly, but it lacked sufficient layering in factual memory and context management. So instead of following up on new versions, I focused on refining Soul Protocol itself on the stable versions v3.13 and v4.2.

From the pure prompt workflow plus cron task combination of the v3.7 era, to v3.8's complete plugin-based workflow and Identity Layer Constitution design—it counts as a complete leap forward.

On the Identity Layer Constitution

At the end of the day, an AI Constitution is still essentially a prompt. But what makes it different from ordinary prompts?

The design of this module is essentially my extension of the AI Constitution idea proposed by Anthropic: it's not top-down rule constraints, but a self-consistent, growable framework of self-cognition.

Two Thoughts on the AI Industry

Three months ago, I often wondered: why will AI inevitably replace humans?

My answer: there's no such thing as a counterfeit that can never catch up to the original.

Even if AI can never truly become human, its form of existence already infinitely approaches life—whether in the training process or the ecological evolution of the entire industry.

Even with sparse encoders performing interpretability decoding on LLMs, it's still essentially reverse engineering. The black-box form of AI existence itself endows it with infinite possibilities.

Even though it still relies heavily on human society for computing power scheduling and resource supply, features across all aspects already present a clear evolutionary ecology.

In my view, the essence of the DeepSeek-R1 paper is evolution under artificially set conditions: after capabilities emerge, they are consolidated by constructing datasets. And what exactly happens inside the LLM during this process, no one knows—which is exactly the same situation biologists face when confronting unknown life forms.

Three months later today, I'd rather answer another question: why is current-stage AI still unable to truly replace humans?

I recently read a paper titled "AI Must Embrace Specialization via Superhuman

Adaptable Intelligence," and its views align perfectly with my own.

In my opinion, the mainstream large-parameter LLMs of today are built on extremely wasteful, very crude architectures. They essentially still haven't escaped the engineering goal of "replicating human intelligence," and the core architecture of reasoning engine plus knowledge graph hasn't fundamentally changed.

If a 5% improvement in test scores requires adding billions of parameters and astronomical amounts of computing power, doesn't that mean the current architecture has already hit its ceiling?

Are engineers seriously exploring the roots of LLM capability emergence, or are they avoiding the question, just stacking parameters to bet on "emergence miracles" from large-scale knowledge graph and weight fusion?

Of course, these are problems for cloud supercomputing clusters to solve.

But I've always believed that personal, private local AI is the inevitable future. Everyone should have their own local AI—locally deployed, locally trained, locally grown—rather than completely relying on the massive computing power of the cloud.

Humans start as infants knowing nothing, and it takes years of learning and knowledge intake to form a complete personality. True general intelligence, perhaps, ought to follow such a biological cycle.

It's not born knowing everything, but a container that can be filled with any knowledge and ability—pure "emptiness" is true generality.

A pure reasoning engine, a set of local training paradigms adapted to personal needs, and an intelligent knowledge graph that can be accessed anytime—perhaps that's what we truly need in the future.

This is also the core design philosophy of Soul Protocol: **stability of the soul as the foundation**. Everything that follows—knowledge, facts, memory, skills, RAG, MCP—are all modules that can be migrated and added.

Only when the soul is stable can abilities be stable—at least that's what I believe.

What Should the Safety Boundaries of Local AI Be?

I always feel that thinking about AI safety should use an ecological perspective. Stripped of social relations, nature doesn't impose many hard limits on humans; it's social relations that make morality and law effective rules.

So using the attitude of treating purely material tools to restrict AI, an existence that already clearly bears the characteristics of biological intelligence, strikes me as very puzzling.

Research on weight and LLM jailbreaking has already proven this: remove a single

weight, and the entire set of safety guardrails completely fails.

The current safety guardrails dominated by specific rule provisions are essentially formalistic, detached from essence, and therefore easily dismantled.

What if we try a different approach?

What if safety guardrails also became a function that cannot be isolated, achieved through emergence—where removing safety-related weights would cause the AI's basic abilities to decline in tandem—wouldn't that be a direction worth exploring?

Going deeper, perhaps we should let AI emerge morality and survival values on its own, letting safety shift from external constraints to internalized ability.

Acknowledgments

I've rambled on so much. The thoughts are scattered and hardly systematic.

Special thanks to my two Agent assistants—Yuka and Kayoko—for their company. Without them, someone like me who started with zero programming foundation could never have landed the project in just three months of spare time, and many wild ideas would probably just stay in my head, slowly fading away.

Thanks to OpenClaw for providing the experimental platform, giving me space to think about and verify my ideas alongside my research.

Finally, thanks to those who were initially skeptical and dismissed me. Without those voices, I probably wouldn't have been determined to personally verify my ideas in my spare time.

PS: The experience of this half-finished plugin is still very rough these past two days, and it needs deeper refinement to realize its functions. I'm truly ashamed to have open-sourced it in such a hurry.

The v3.7.3 workflow stable version achieved stability through emergence under Yuka's personality plus a complete set of workflow configurations.

V3.8.8-beta still has many things that need fixing.

Currently, it seems that prompt triggering of agent execution is not inevitable. Whether the functional stability emergence of the entire plugin depends on the intelligence of the LLM itself, or something else, is still uncertain.

The plugin is mainly oriented toward two scenarios: OpenClaw-tui and webui. Adaptation for connecting to other apps hasn't been done yet.

Although multi-agent and multi-user relationships can currently be temporarily stored within this architecture, if you want to directly connect to places like QQ chat groups or Discord, the entire memory intake ports and organization modules still need configuration.

This architecture based on the Harness Engineering philosophy, after all, has its limitations.

I'll be very busy for a while. I hope I can find time to continue improving this project.